



Algorithms: Design
and Analysis, Part II

Greedy Algorithms

A Scheduling Application:
Problem Definition

A Scheduling Problem

Setup:

- One shared resource (e.g., a processor).
- Many “jobs” to do (e.g., processes).

Question: In what order should we sequence the jobs?

Assume: Each job has a:

- weight w_j (“priority”)
- length l_j

Completion Times

Definition: The completion time C_j of job j = Sum of job lengths up to and including j .

Example: 3 jobs, $l_1 = 1, l_2 = 2, l_3 = 3$.

Schedule:

#1	#2	#3
----	----	----

0 →

(time)

Question: What is C_1, C_2, C_3 ?

A) 1, 2, 3

C) 1, 3, 6

B) 3, 5, 6

D) 1, 4, 6

The Objective Function

Goal: Minimize the weighted sum of completion times:

$$\min \sum_{j=1}^n w_j C_j.$$

Back to example: If $w_1 = 3$, $w_2 = 2$, $w_3 = 1$, this sum is $3 \cdot 1 + 2 \cdot 3 + 1 \cdot 6 = 15$.



Algorithms: Design
and Analysis, Part II

Greedy Algorithms

A Scheduling Application:
The Algorithm

Intuition for Algorithm

Recall: Want to $\min \sum_{j=1}^n w_j \cdot C_j$

Goal: Devise correct greedy algorithm.

Question:

1. With equal lengths, schedule larger or smaller-weight jobs earlier?
2. With equal weights, schedule shorter or longer jobs earlier?

A) Larger/shorter

C) Larger/longer

B) Smaller/shorter

D) Smaller/longer

Resolving Conflicting Advice

Question: What if $w_i > w_j$ but $l_i > l_j$?

Idea: Assign “scores” to jobs that are:

- inscreasing in weight
- decreasing in length

Guess (1): Order jobs by decreasing value of $w_j - l_j$.

Guess (2): Order w_j/l_j .

Breaking a Greedy Algorithm

To distinguish (1) & (2): Find example where the two algorithms produce different outputs. (At least one will be incorrect.)

Example:

$$l_1 = 5, w_1 = 3 \text{ (longer ratio)} \quad w1 = 3, l1 = 5$$

$$l_1 = 2, w_1 = 1 \text{ (larger difference)} \quad w2 = 1, l2 = 2$$

Question: What is the sum of weighted completion times of algorithms (1) & (2) respectively?

A) 22 and 23 C) 17 and 17

B) 23 and 22 D) 17 and 11

$$\text{Alg\#1: } \begin{array}{|c|c|} \hline \#2 & \#1 \\ \hline \end{array} \rightarrow 1 \cdot 2 + 3 \cdot 7 = 23$$

$$\text{Alg\#2: } \begin{array}{|c|c|} \hline \#1 & \#2 \\ \hline \end{array} \rightarrow 3 \cdot 5 + 1 \cdot 7 = 22$$

The Story So Far

So: Alg#1 not (always) correct.

Claim: Alg#2 (order by decreasing ratio w_j/l_j 's) is always correct.
[not obvious! - proof coming up next]

Running time: $O(n \log n)$. [just need to sort]



Algorithms: Design
and Analysis, Part II

Greedy Algorithms

A Scheduling Application:
Correctness Proof Part I

Correctness Claim

Claim: Algorithm #2 (order jobs according to decreasing ratios w_j/l_j) is always correct.

Proof: By an **Exchange Argument**.

Plan: Fix arbitrary input of n jobs. Will proceed by contradiction. Let σ = greedy schedule, σ^* = optimal schedule. (With σ^* better than σ .)

Will produce schedule even better than σ^* , contradicting purported optimality of σ^* .

Correctness Proof

Assume: All w_j/l_j 's distinct.

Assume: [Just by renaming jobs] $w_1/l_1 > w_2/l_2 > \dots > w_n/l_n$.

Thus: Greedy schedule σ is just $1, 2, 3, \dots, n$.

Thus: If optimal schedule $\sigma^* \neq \sigma$, then there are consecutive jobs i, j with $i > j$.

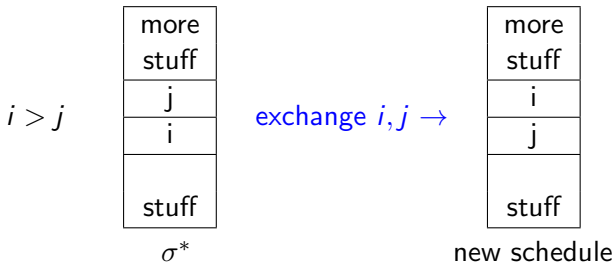
[Only schedule where indices always go up is $1, 2, 3, \dots, n$]

Correctness Proof (con'd)

So far:

1. $w_1/l_1 > w_2/l_2 > \dots > w_n/l_n$
2. In optimal σ^* , \exists consecutive jobs i, j with $i > j$.

Thought experiment: Suppose we **exchange** order of i & j in σ^* (leaving other jobs unchanged):






Algorithms: Design
and Analysis, Part II

Greedy Algorithms

A Scheduling Application:
Correctness Proof Part II

Cost-Benefit Analysis, Part I



Question: What is the effect of this exchange on the completion time of (1) a job k other than i or j , (2) the job i , (3) the job j ?

- A) Not enough info/goes up/goes down
 - B) Not enough info/goes down/goes up
 - C) Unaffected/ goes up / goes down
 - D) Unaffected/goes down/goes up
- by j
- by i

Cost-Benefit Analysis, Part II

Upshot:

1. Cost of exchange $w_i l_j$. [C_i goes up by l_j]
2. Benefit of exchange is $w_j l_i$. [C_j goes down by l_i]

Note: $i > j \Rightarrow w_i/l_i < w_j/l_j \Rightarrow w_i l_j < w_j l_i \Rightarrow \text{COST} < \text{BENEFIT}$
 \Rightarrow Swap improves σ^* , contradicts optimality of σ^* .

QED!



Algorithms: Design
and Analysis, Part II

Greedy Algorithms

A Scheduling Application:
Handling Ties

Correctness Claim

Claim: Algorithm #2 (order jobs in nonincreasing order of ratio w_j/l_j) is always correct. [Even with ties]

New Proof Plan: Fix arbitrary input of n jobs. Let σ = greedy schedule, let σ^* = any other schedule.

Will show σ at least as good as $\sigma^* \Rightarrow$ Implies that greedy schedule is optimal.

Correctness Proof


Assume: [Just by renaming jobs] Greedy schedule σ is just $1, 2, 3, \dots, n$ (and so $w_1/l_1 \geq w_2/l_2 \geq \dots \geq w_n/l_n$).

Consider arbitrary schedule σ^* . If $\sigma^* = \sigma$, done.

Else recall \exists consecutive jobs i, j in σ^* with $i > j$. (From last time)

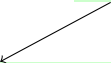
Note: $i > j \Rightarrow w_i/l_i \leq w_j/l_j \Rightarrow w_i l_j \leq w_j l_i$.

Recall: Exchanging i & j in σ^* has net benefit of $w_j l_i - w_i l_j \geq 0$.



Correctness Proof

Upshot: Exchanging an “adjacent inversion” like i, j only makes σ^* better, and it decreases the number of **inverted pairs**.



Jobs i, j with $i > j$ and i scheduled earlier

\Rightarrow After at most $\binom{n}{2}$ such exchanges, can transform σ^* into σ .

$\Rightarrow \sigma$ at least as good as σ^* .

\Rightarrow Greedy is optimal.

QED!