Minimum Spanning Trees

Problem Definition

Algorithms: Design and Analysis, Part II

Overview

Informal Goal: Connect a bunch of points together as cheaply as possible.

Applications: Clustering (more later), networking.

Blazingly Fast Greedy Algorithms:

- Prim's Algorithm [1957; also Dijkstra 1959, Jarnik 1930]
- Kruskal's algorithm [1956]

$\Rightarrow O(\boxed{m} \log \boxed{n})$ time (using suitable data structures)

# of vertices

# of edges

Problem Definition

vertices　　　　　edges

Input: Undirected graph $G = (V, E)$ and a cost c_e for each edge $e \in E$.

- Assume adjacency list representation (see Part I for details)
- OK if edge costs are negative

Output: minimum cost tree $T \subseteq E$ that spans all vertices.

i.e., sum of edge costs

I.e.: (1) T has no cycles, (2) the subgraph (V, T) is connected (i.e., contains path between each pair of vertices).

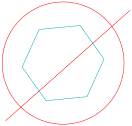

(disallowed)

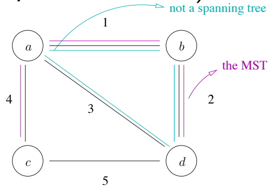

Standing Assumptions

Assumption #1: Input graph G is connected.

- Else no spanning trees.
- Easy to check in preprocessing (e.g., depth-first search).

Assumption #2: Edge costs are distinct.

- Prim + Kruskal remain correct with ties (which can be broken arbitrarily).
- Correctness proof a bit more annoying (will skip).

Minimum Spanning Trees

Prim's MST Algorithm

Algorithms: Design and Analysis, Part II

Example

[Purple edges = minimum spanning tree]

(Compare to Dijkstra's shortest-path algorithm)

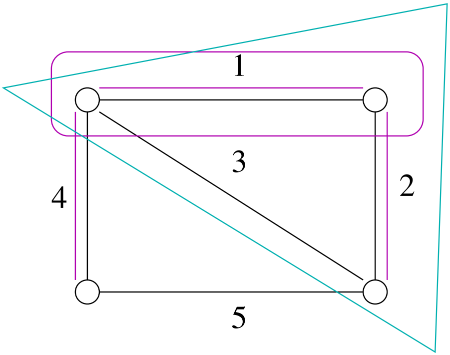

Prim's MST Algorithm

- Initialize $X = \{s\}$ [$s \in V$ chosen arbitrarily]
- $T = \emptyset$ [invariant: X = vertices spanned by tree-so-far T]
- While $X \neq V$
  - Let $e = (u, v)$ be the cheapest edge of G with $u \in X$, $v \notin X$.
  - Add e to T
  - Add v to X.

While loop: Increase # of spanned vertices in cheapest way possible.

Correctness of Prim's Algorithm

Theorem: Prim's algorithm always computes an MST.

Part I: Computes a spanning tree T^*.
[Will use basic properties of graphs and spanning trees] (Useful also in Kruskal's MST algorithm)

Part II: T^* is an MST.
[Will use the "Cut Property"] (Useful also in Kruskal's MST algorithm)

Later: Fast [$O(m \log n)$] implementation using heaps.

Minimum Spanning Trees

Algorithms: Design and Analysis, Part II

Correctness of Prim's Algorithm (Part I)

Cuts

Claim: Prim's algorithm outputs a spanning tree.

Definition: A cut of a graph $G = (V, E)$ is a partition of V into 2 non-empty sets.

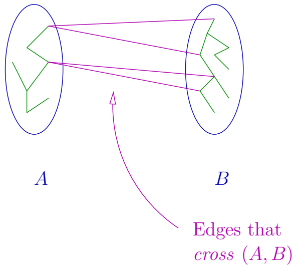

Quiz on Cuts

Question: Roughly how many cuts does a graph with n vertices have?

A) n C) 2^n (for each vertex, choose whether in A or in B)

B) n^2 D) n^n

Empty Cut Lemma

Empty Cut Lemma: A graph is not connected $\iff \exists$ cut (A, B) with no crossing edges.

Proof: ($\Leftarrow$) Assume the RHS. Pick any $u \in A$ and $v \in B$. Since no edges cross (A, B) there is no u, v path in G. $\Rightarrow G$ not connected.

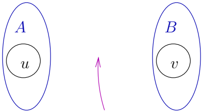

($\Rightarrow$) Assume the LHS. Suppose G has no $u - v$ path. Define

$A = \{$Vertices reachable from u in $G\}$ (u's connected component)

$B = \{$All other vertices$\}$ (all other connected components)

Note: No edges cross cut (A, B) (otherwise A would be bigger!)

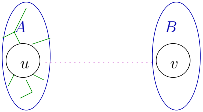

QED!

Two Easy Facts

Double-Crossing Lemma: Suppose the cycle $C \subseteq E$ has an edge crossing the cut (A, B): then so does some other edge of C.

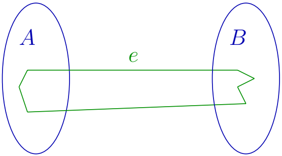

Lonely Cut Corollary: If e is the only edge crossing some cut (A, B), then it is not in any cycle. [If it were in a cycle, some other edge would have to cross the cut!]

Proof of Part I

Claim: Prim's algorithm outputs a spanning tree.
[Not claiming MST yet]

Proof: (1) Algorithm maintains invariant that T spans X
[straightforward induction - you check]

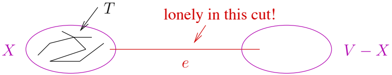

(2) Can't get stuck with $X \neq V$
[otherwise the cut $(X, V - X)$ must be empty; by Empty Cut Lemma input graph G is disconnected]

(3) No cycles ever get created in T. Why? Consider any iteration, with current sets X and T. Suppose e gets added.
Key point: e is the first edge crossing $(X, V - X)$ that gets added to $T \Rightarrow$ its addition can't create a cycle in T (by Lonely Cut Corollary). QED!

Minimum Spanning Trees

Algorithms: Design and Analysis, Part II

Correctness of Prim's Algorithm (Part II)

Correctness of Prim's Algorithm

Theorem: Prim's algorithm always outputs a minimum-cost spanning tree.

Key Question: When is it "safe" to include an edge in the tree-so-far?

The Cut Property

CUT PROPERTY: Consider an edge e of G. Suppose there is a cut (A, B) such that e is the cheapest edge of G that crosses it. Then e belongs to the MST of G.

Turns out MST is unique if edge costs are distinct

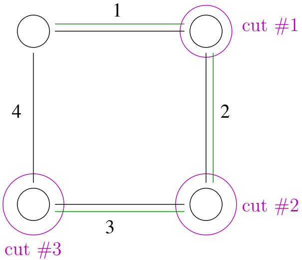

Cut Property Implies Correctness

Claim: Cut Property $\Rightarrow$ Prim's algorithm is correct.

Proof: By previous video, Prim's algorithm outputs a spanning tree T^*.

Key point: Every edge $e \in T^*$ is explicitly justified by the Cut Property.

$\Rightarrow$ T^* is a subset of the MST

$\Rightarrow$ Since T^* is already a spanning tree, it must be the MST

QED!

Minimum Spanning Trees

Proof of the Cut Property

Algorithms: Design and Analysis, Part II

The Cut Property

Assumption: Distinct edge costs.

CUT PROPERTY: Consider an edge e of G. Suppose there is a cut (A, B) such that e is the cheapest edge of G that crosses it. Then e belongs to the MST of G.

Proof Plan

Will argue by contradiction, using an exchange argument.
[Compare to scheduling application]

Suppose there is an edge e that is the cheapest one crossing a cut (A, B), yet e is not in the MST T^*.

Idea: Exchange e with another edge in T^* to make it even cheaper (contradiction).

Question: Which edge to exchange e with?

Attempted Exchange

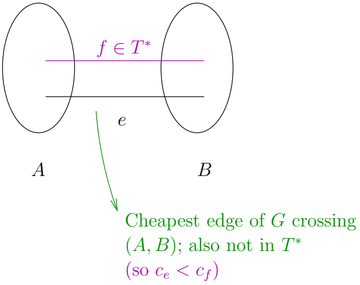

Note: Since T^* is connected, must construct an edge $f (\neq e)$ crossing (A, B).

Idea: Exchange e and f to get a spanning tree cheaper than T^* (contradiction).

Exchanging Edges

Question: Let T^* be a spanning tree of G, $e \notin T^*$, $f \in T^*$. Is $T^* \cup \{e\} - \{f\}$ a spanning tree of G?

A) Yes always
B) No never
C) If e is the cheapest edge crossing some cut, then yes
D) Maybe, maybe not (depending on the choice of e and f)

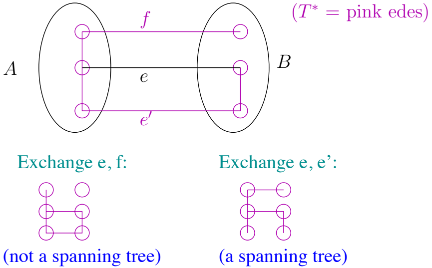

Smart Exchanges

Hope: Can always find suitable edge e' so that exchange yields bona fide spanning tree of G.

How? Let C = cycle created by adding e to T^*.

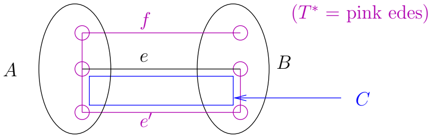

By the Double-Crossing Lemma: Some other edge e' of C [with $e' \neq e$ and $e' \in T^*$] crosses (A, B).

You check: $T = T^* \cup \{e\} - \{e'\}$ is also a spanning tree.

Since $c_e < c_{e'}$, T cheaper than purported MST T^*, contradiction.

Tim Roughgarden

Minimum Spanning Trees

Fast Implementation of Prim's Algorithm

Algorithms: Design and Analysis, Part II

Running Time of Prim's Algorithm

- Initialize $X = \{s\}$ [$s \in V$ chosen arbitrarily]
- $T = \emptyset$ [invariant: X = vertices spanned by tree-so-far T]
- While $X \neq V$
    - Let $e = (u, v)$ be the cheapest edge of G with $u \in X$, $v \notin X$.
    - Add e to T, add v to X.

Running time of straightforward implementation:
- $O(n)$ iterations [where n = # of vertices]
- $O(m)$ time per iteration [where m = # of edges]

$\Rightarrow O(mn)$ time

BUT CAN WE DO BETTER?

Prim's Algorithm with Heaps

[Compare to fast implementation of Dijkstra's algorithm]

Invariant #1: Elements in heap = vertices of $V - X$.

Invariant #2: For $v \in V - X$, key[v] = cheapest edge (u, v) with $i \in X$ (or $+\infty$ if no such edges exist).

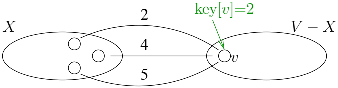

Check: Can initialize heap with $O(m + n \log n) = O(m \log n)$ preprocessing.

To compare keys | $n - 1$ Inserts | $m \geq n - 1$ since G connected

Note: Given invariants, Extract-Min yields next vertex $v \notin X$ and edge (u, v) crossing $(X, V - X)$ to add to X and T, respectively.

Quiz: Issue with Invariant #2

Question: What is: (i) current value of key[v] (ii) current value of key[w] (iii) value of key[w] after one more iteration of Prim's algorithm?

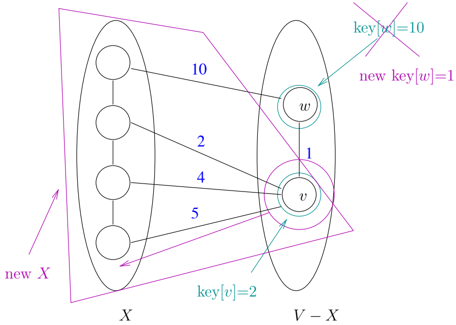

A) 11, 10, 4

B) 2, 10, 10

C) 2, 10, 1

D) 2, 10, 2

Maintaining Invariant #2

Issue: Might need to recompute some keys to maintain Invariant #2 after each Extract-Min.

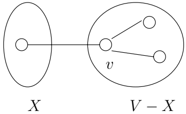

Pseudocode: When v added to X:

- For each edge $(v, w) \in E$:
  - If $w \in V - X$ $\rightarrow$ The only whose key might have changed (Update key if needed:)
    - Delete w from heap
    - Recompute key[w]:=min{key[w],c_{vw}}
    - Re-Insert into heap

Subtle point/exercise:

Think through book-keeping needed to pull this off

Running Time with Heaps

- Dominated by time required for heap operations
- $(n-1)$ Inserts during preprocessing
- $(n-1)$ Extract-Mins (one per iteration of while loop)
- Each edge (v, w) triggers one Delete/Insert combo
  [When its first endpoint is sucked into X]

$\Rightarrow$ $O(m)$ heap operations [Recall $m \geq n-1$ since G connected]

$\Rightarrow$ $O(m \log n)$ time [As fast as sorting!]